
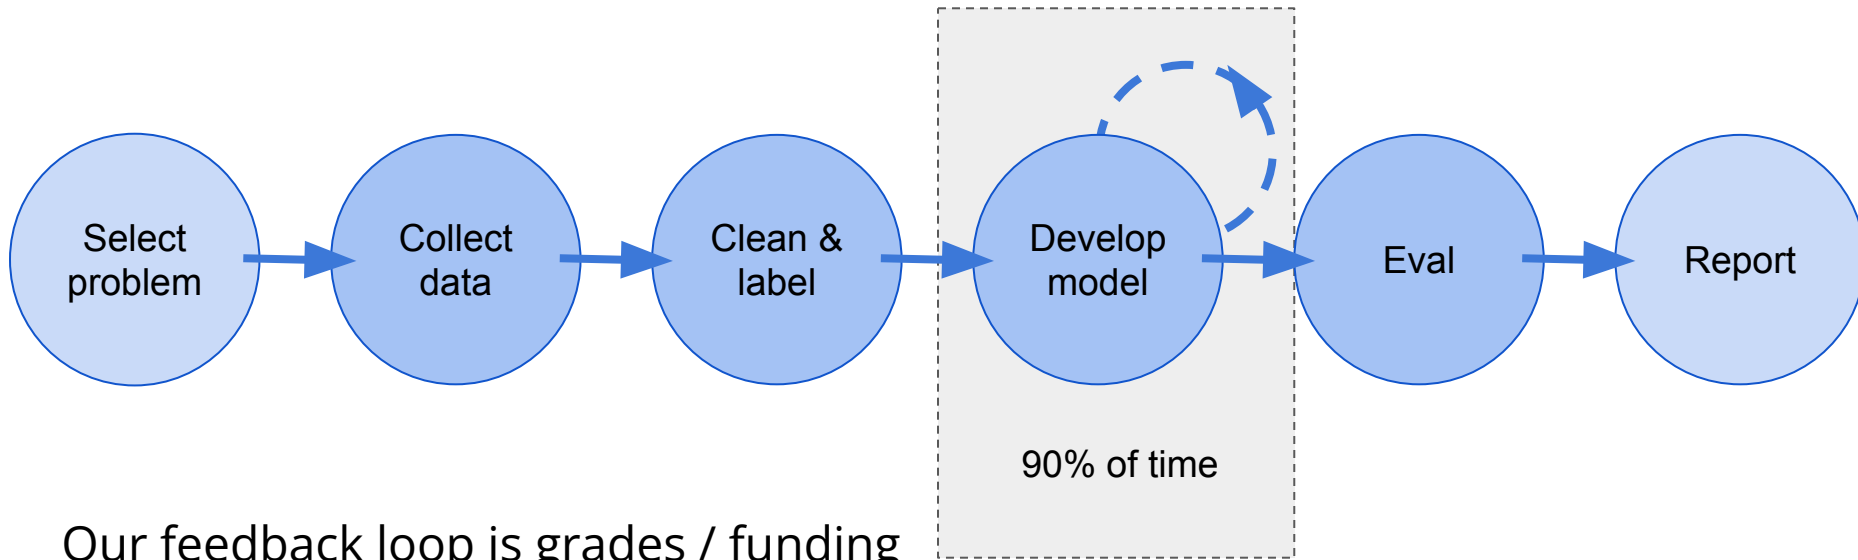

What is Machine Learning Operations?

— 02476 Machine Learning Operations —
Nicki Skafte Detlefsen



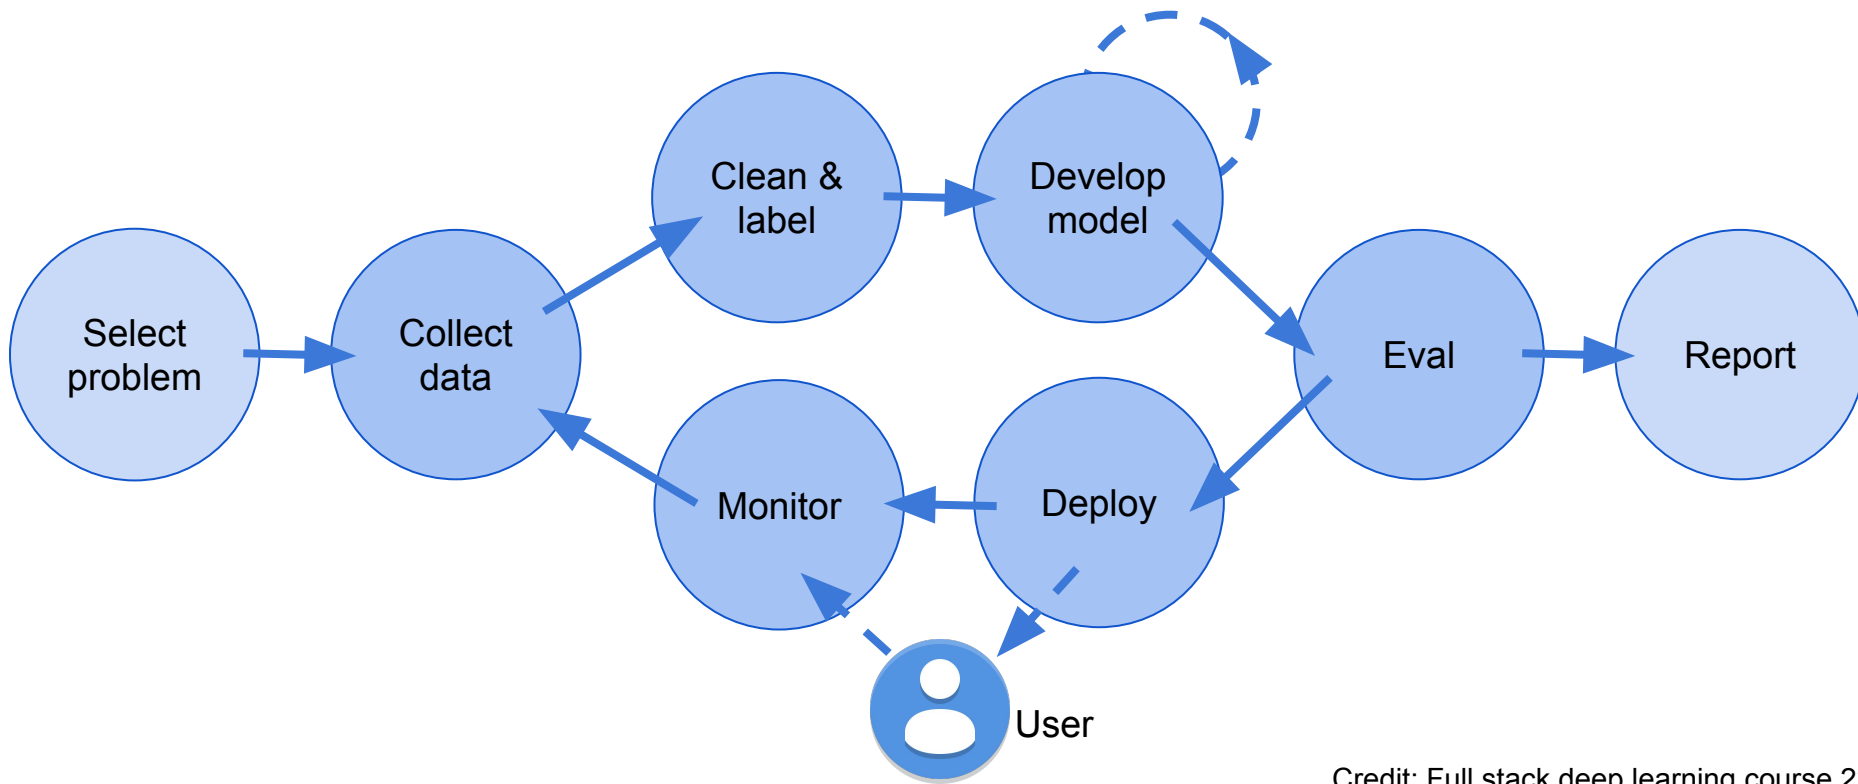

Let start where you are now

Courses / Projects are linear in nature

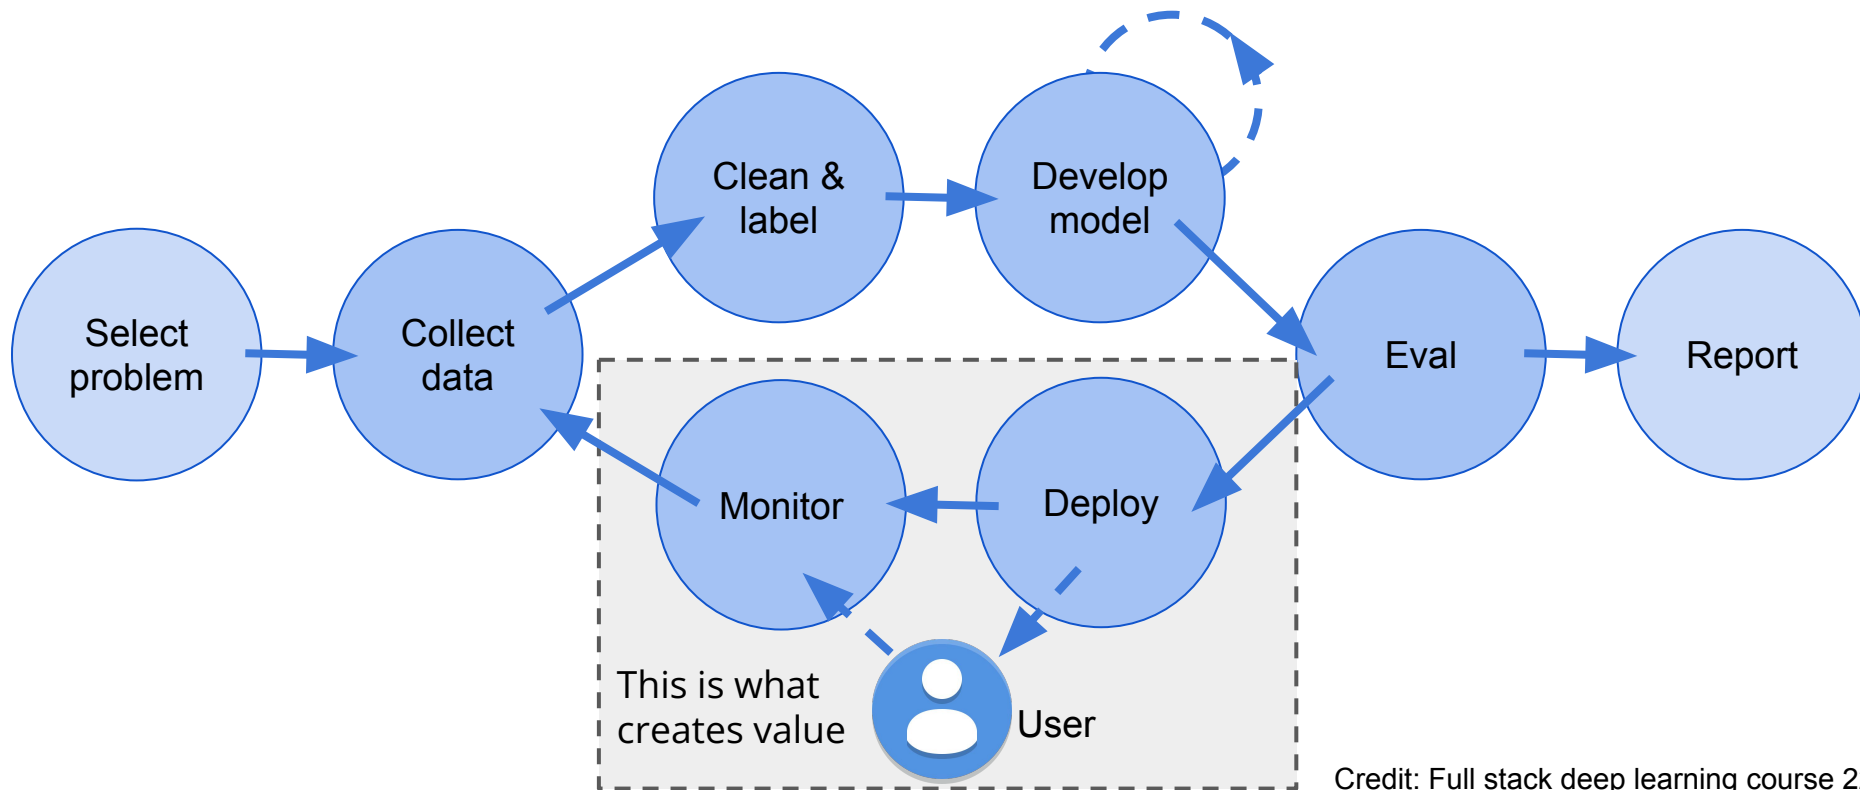



Our feedback loop is grades / funding

ML in the real world

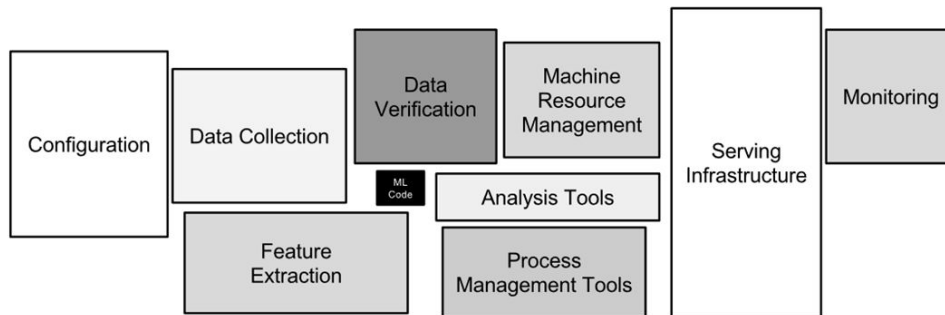


ML in the real world




Key observations

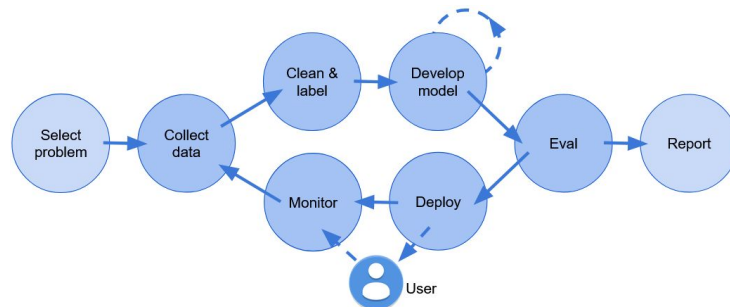
1. Machine learning in production is much more than doing ML modelling



2. Machine Learning in production is a cycle



VS.




What is MLOps?

This is both a joke and not.

MLOps is directly derived from DevOps.

Therefore, let's try to understand DevOps first.



So, what is DevOps?


This is the closest to a definition that I could find:

DevOps is a set of practices that combines software development (*Dev*) and IT operations (*Ops*). It aims to shorten the systems development life cycle and provide continuous delivery with high software quality. It's an combination of human mindset, processes and technologies that continuously creates value.

So, what is DevOps?


This is the closest to a definition that I could find:

DevOps is a set of practices that combines software development (*Dev*) and IT operations (*Ops*). It aims to shorten the systems development **life cycle** and provide continuous delivery with high software quality. It's an combination of human **mindset, processes and technologies** that continuously **creates value**.





But then MLOps must be...


Is a set of tools, processes, and mindset that aim to make **ML Lifecycle** reproducible, trackable, testable and maintainable to continuously create value.



Design


- Business understanding
- Data understanding
- Designing the ML-powered software






Model development


- Model engineering
- Data engineering
- Deliver a stable quality ML model that we will run in production




Operations



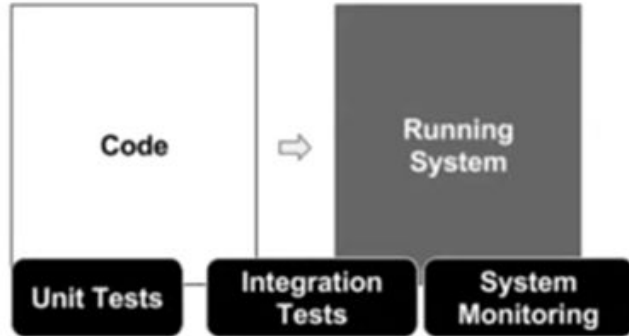
- Deliver the previously developed ML model in production
- Testing, versioning, continuous delivery, and monitoring



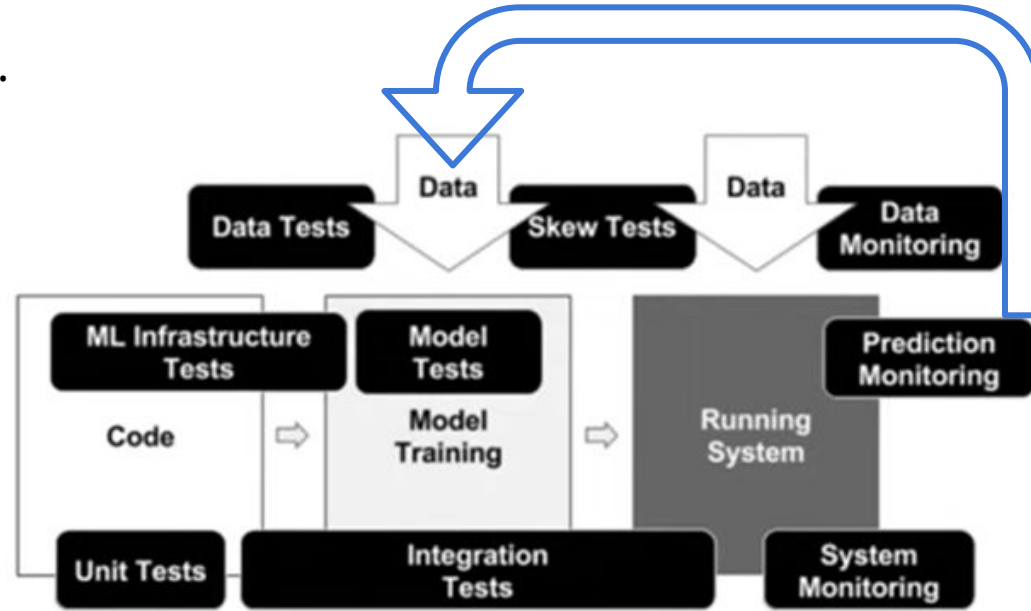


Why is DevOps not enough?


Because *data* changes everything...



Traditional System Testing and Monitoring



ML-Based System Testing and Monitoring




Why does companies care about MLOps

Having automated model deployed with errors can cost A LOT of money:

"A famous example of the dangers here was Knight Capital's system losing \$465 millions in 45 minutes, apparently because of unexpected behavior from obsolete experimental codepaths" – Hidden Technical depth in Machine Learning Systems


The Impact of AI Automation on Jobs



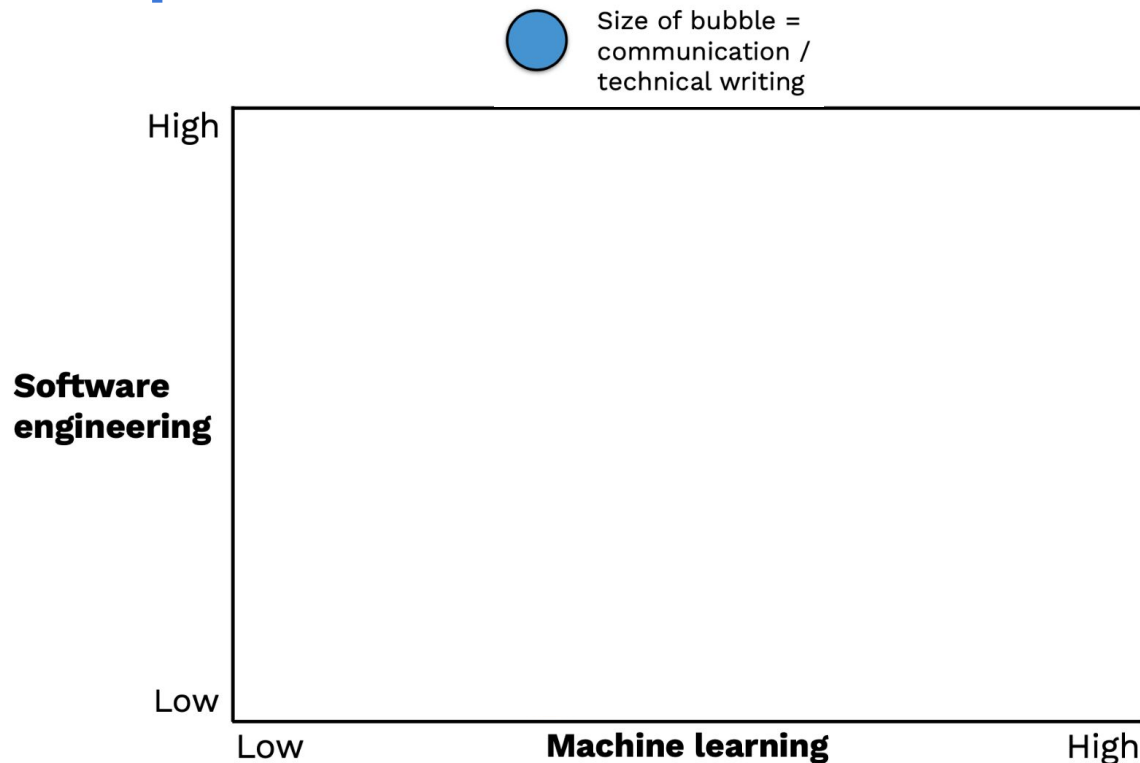
What makes an MLOps engineer?

An mix of


- Software developing
- Machine Learning
- Data Engineering




MLOps vs ML Researcher vs Data scientist

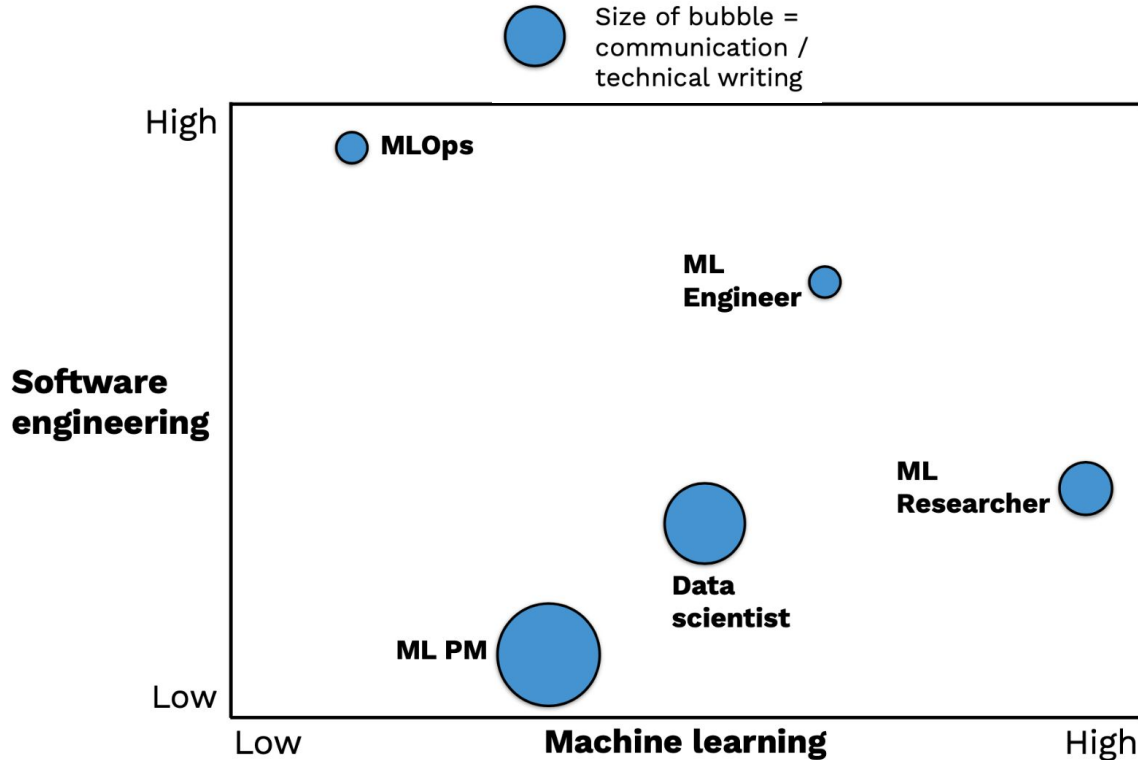


Where should different positions be?





MLOps vs ML Researcher vs Data scientist




Embrace that you cannot do everything

MLOps trends

MLOps has been trending for a couple of years.


Tools have been the main priority

Main priority




Tools

+



Processes

+



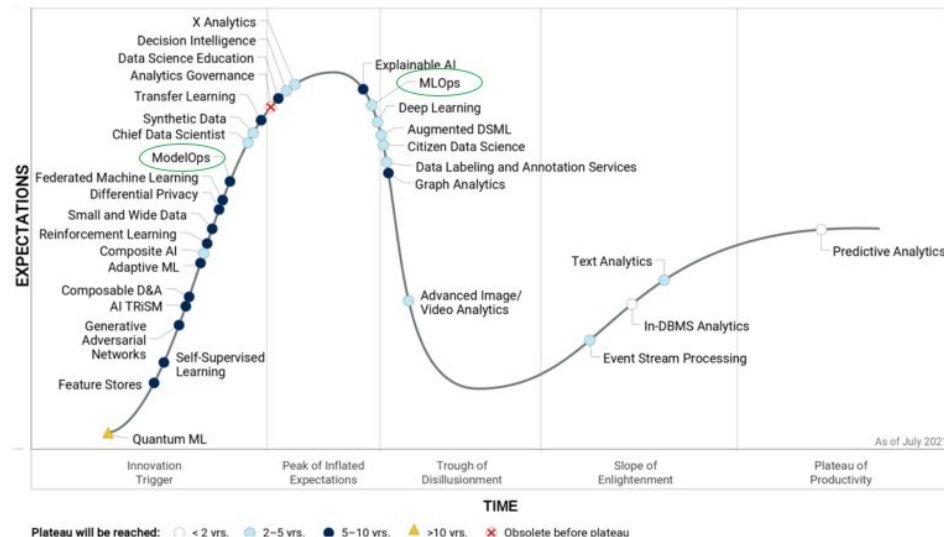
Mindset

=



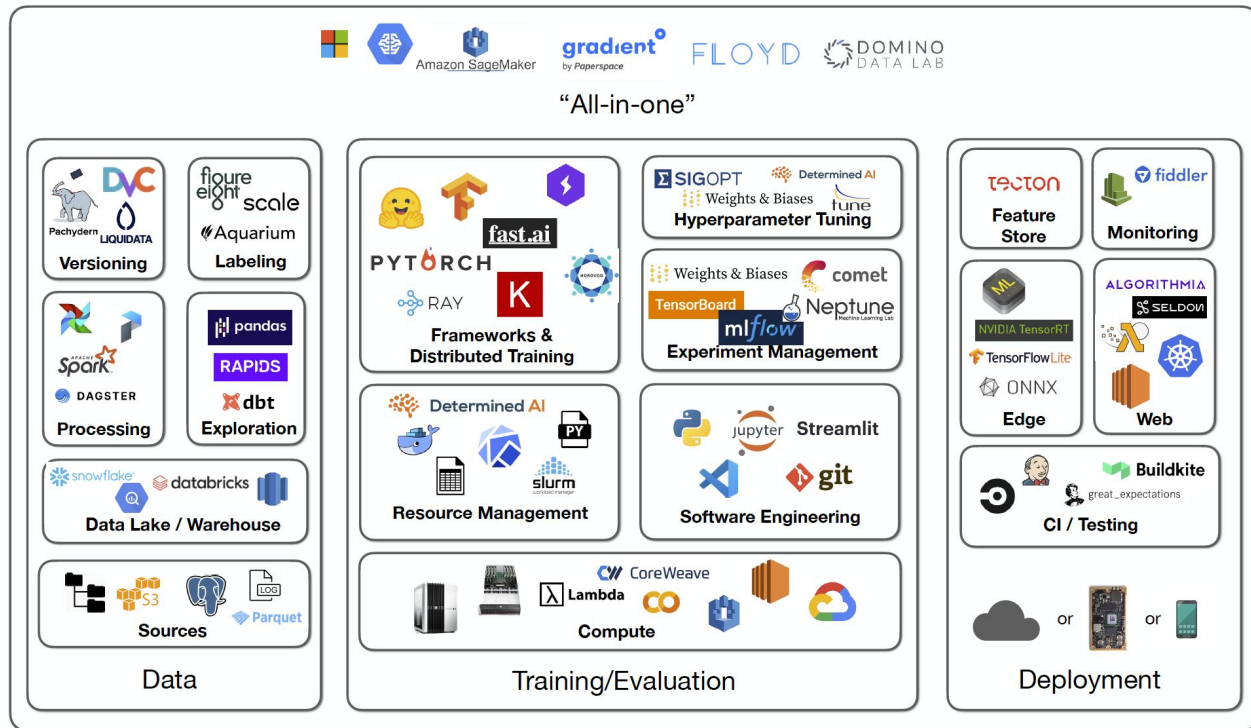
Value

Credit: Gartner 2021 Hype report



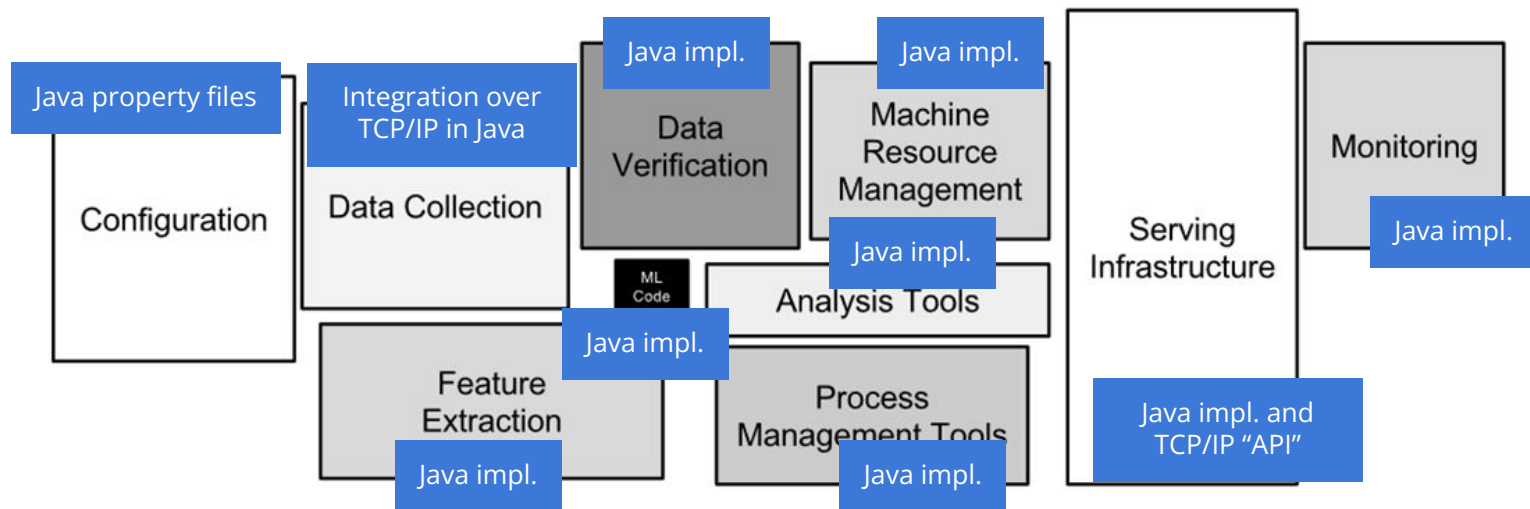
MLOps trends: TOOLS!

Anything you need
there is a tool for



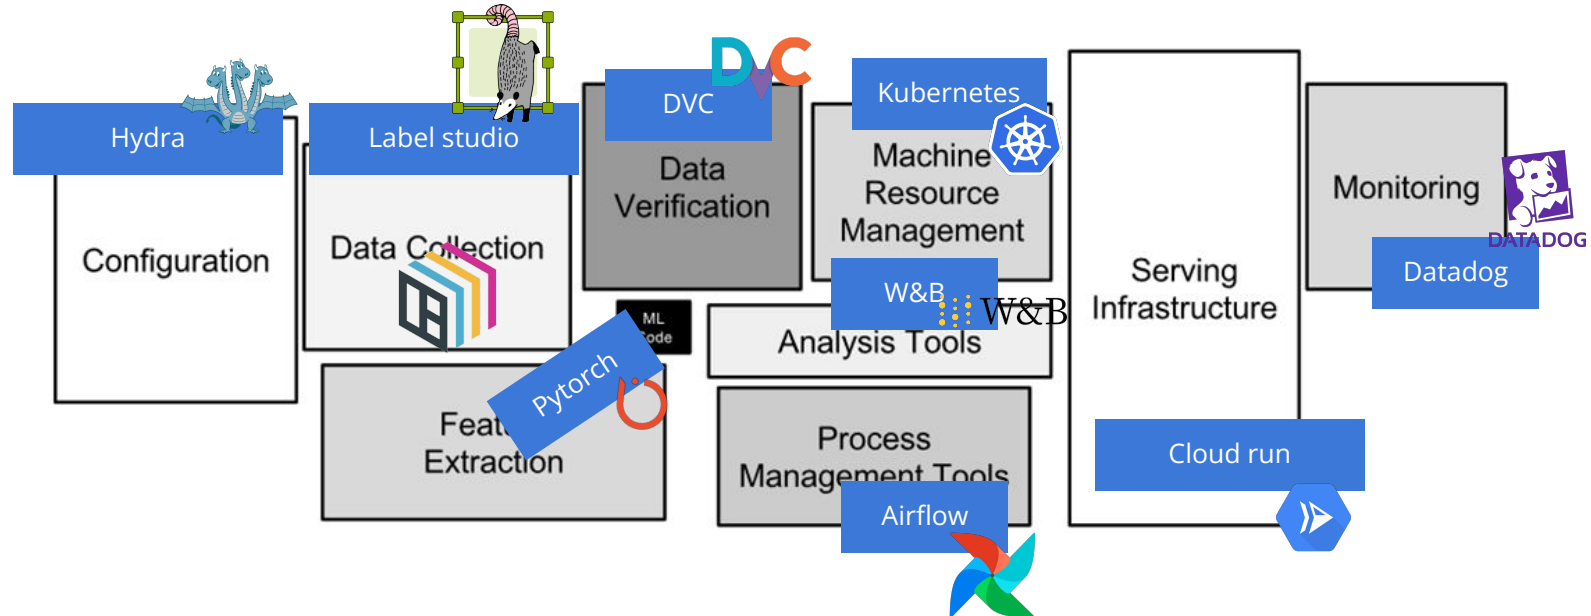
MLOps in the past

MLOps around 2006 = write everything from scratch



MLOps now

Pick a *stack* of tools




MLOps is kind of full stack engineering

In MLOps we embrace the full stack of problems that comes from the full lifecycle.

Criteria for what goes into the stack:

- Cost
- Flexibility
- Complexity





The core challenges in MLOps

1. Deployment and automation
2. Reproducibility of models and predictions
3. Diagnostics
4. Governance and regulatory compliance
5. Scalability
6. Collaboration
7. Business uses
8. Monitoring and management

Meme of the day

